


From GPUs to RRAMs: Distributed In-Memory Primal–Dual Hybrid Gradient Method for Solving Large-Scale Linear Optimization Problems

Huynh Q. N. Vo^{1,2} M. T. R. Chowdhury^{3,2} Paritosh Ramanan¹ Gözde Tutuncuoglu³ Junchi Yang² Feng Qiu² Murat Yildirim⁴

¹Oklahoma State University (IEM) • ²Argonne National Laboratory (ESIA) • ³Wayne State University (ECE) • ⁴Wayne State University (ISE)



ABSTRACT

Large-scale linear optimization underpins decision-making across science, engineering, and industry, yet its growing computational demand increasingly exceeds the efficiency of **conventional von Neumann hardware**.


We present the **first primal-dual hybrid gradient (PDHG) linear-program solver co-designed for in-memory computing on resistive-RAM (RRAM) crossbar arrays**. Evaluated in a physics-based device simulation, the solver attains **accuracy comparable to a GPU** while reducing energy and latency by up to **three orders of magnitude**.

MOTIVATION — Where linear programs arise

A linear program (LP) optimizes a linear **objective function** (e.g. cost, energy, or profit) subject to linear **constraints** that encode real-world limits (budgets, capacities, supply and demand):

$$\min_{\mathbf{x} \in \mathbb{R}^n} \mathbf{c}^\top \mathbf{x}, \quad \text{s.t.} \quad \mathbf{A}\mathbf{x} \leq \mathbf{b}; \quad l \leq x_i \leq u, \quad \forall x_i \in \mathbf{x}$$


Such problems underpin optimal decision-making across modern infrastructure:



MOTIVATION — The memory wall and in-memory computing

Conventional von Neumann architectures separate the **processor** from the **memory**. For the large matrices underlying these problems, most energy and time is consumed *transferring data* between the two — the **memory wall**.

In-memory computing via RRAM technology eliminates this overhead: RRAM cells encode numerical values as electrical conductances arranged in a grid (a *crossbar*); thus, applying input voltages yields a complete matrix-vector product as output currents in a **single analog step**.

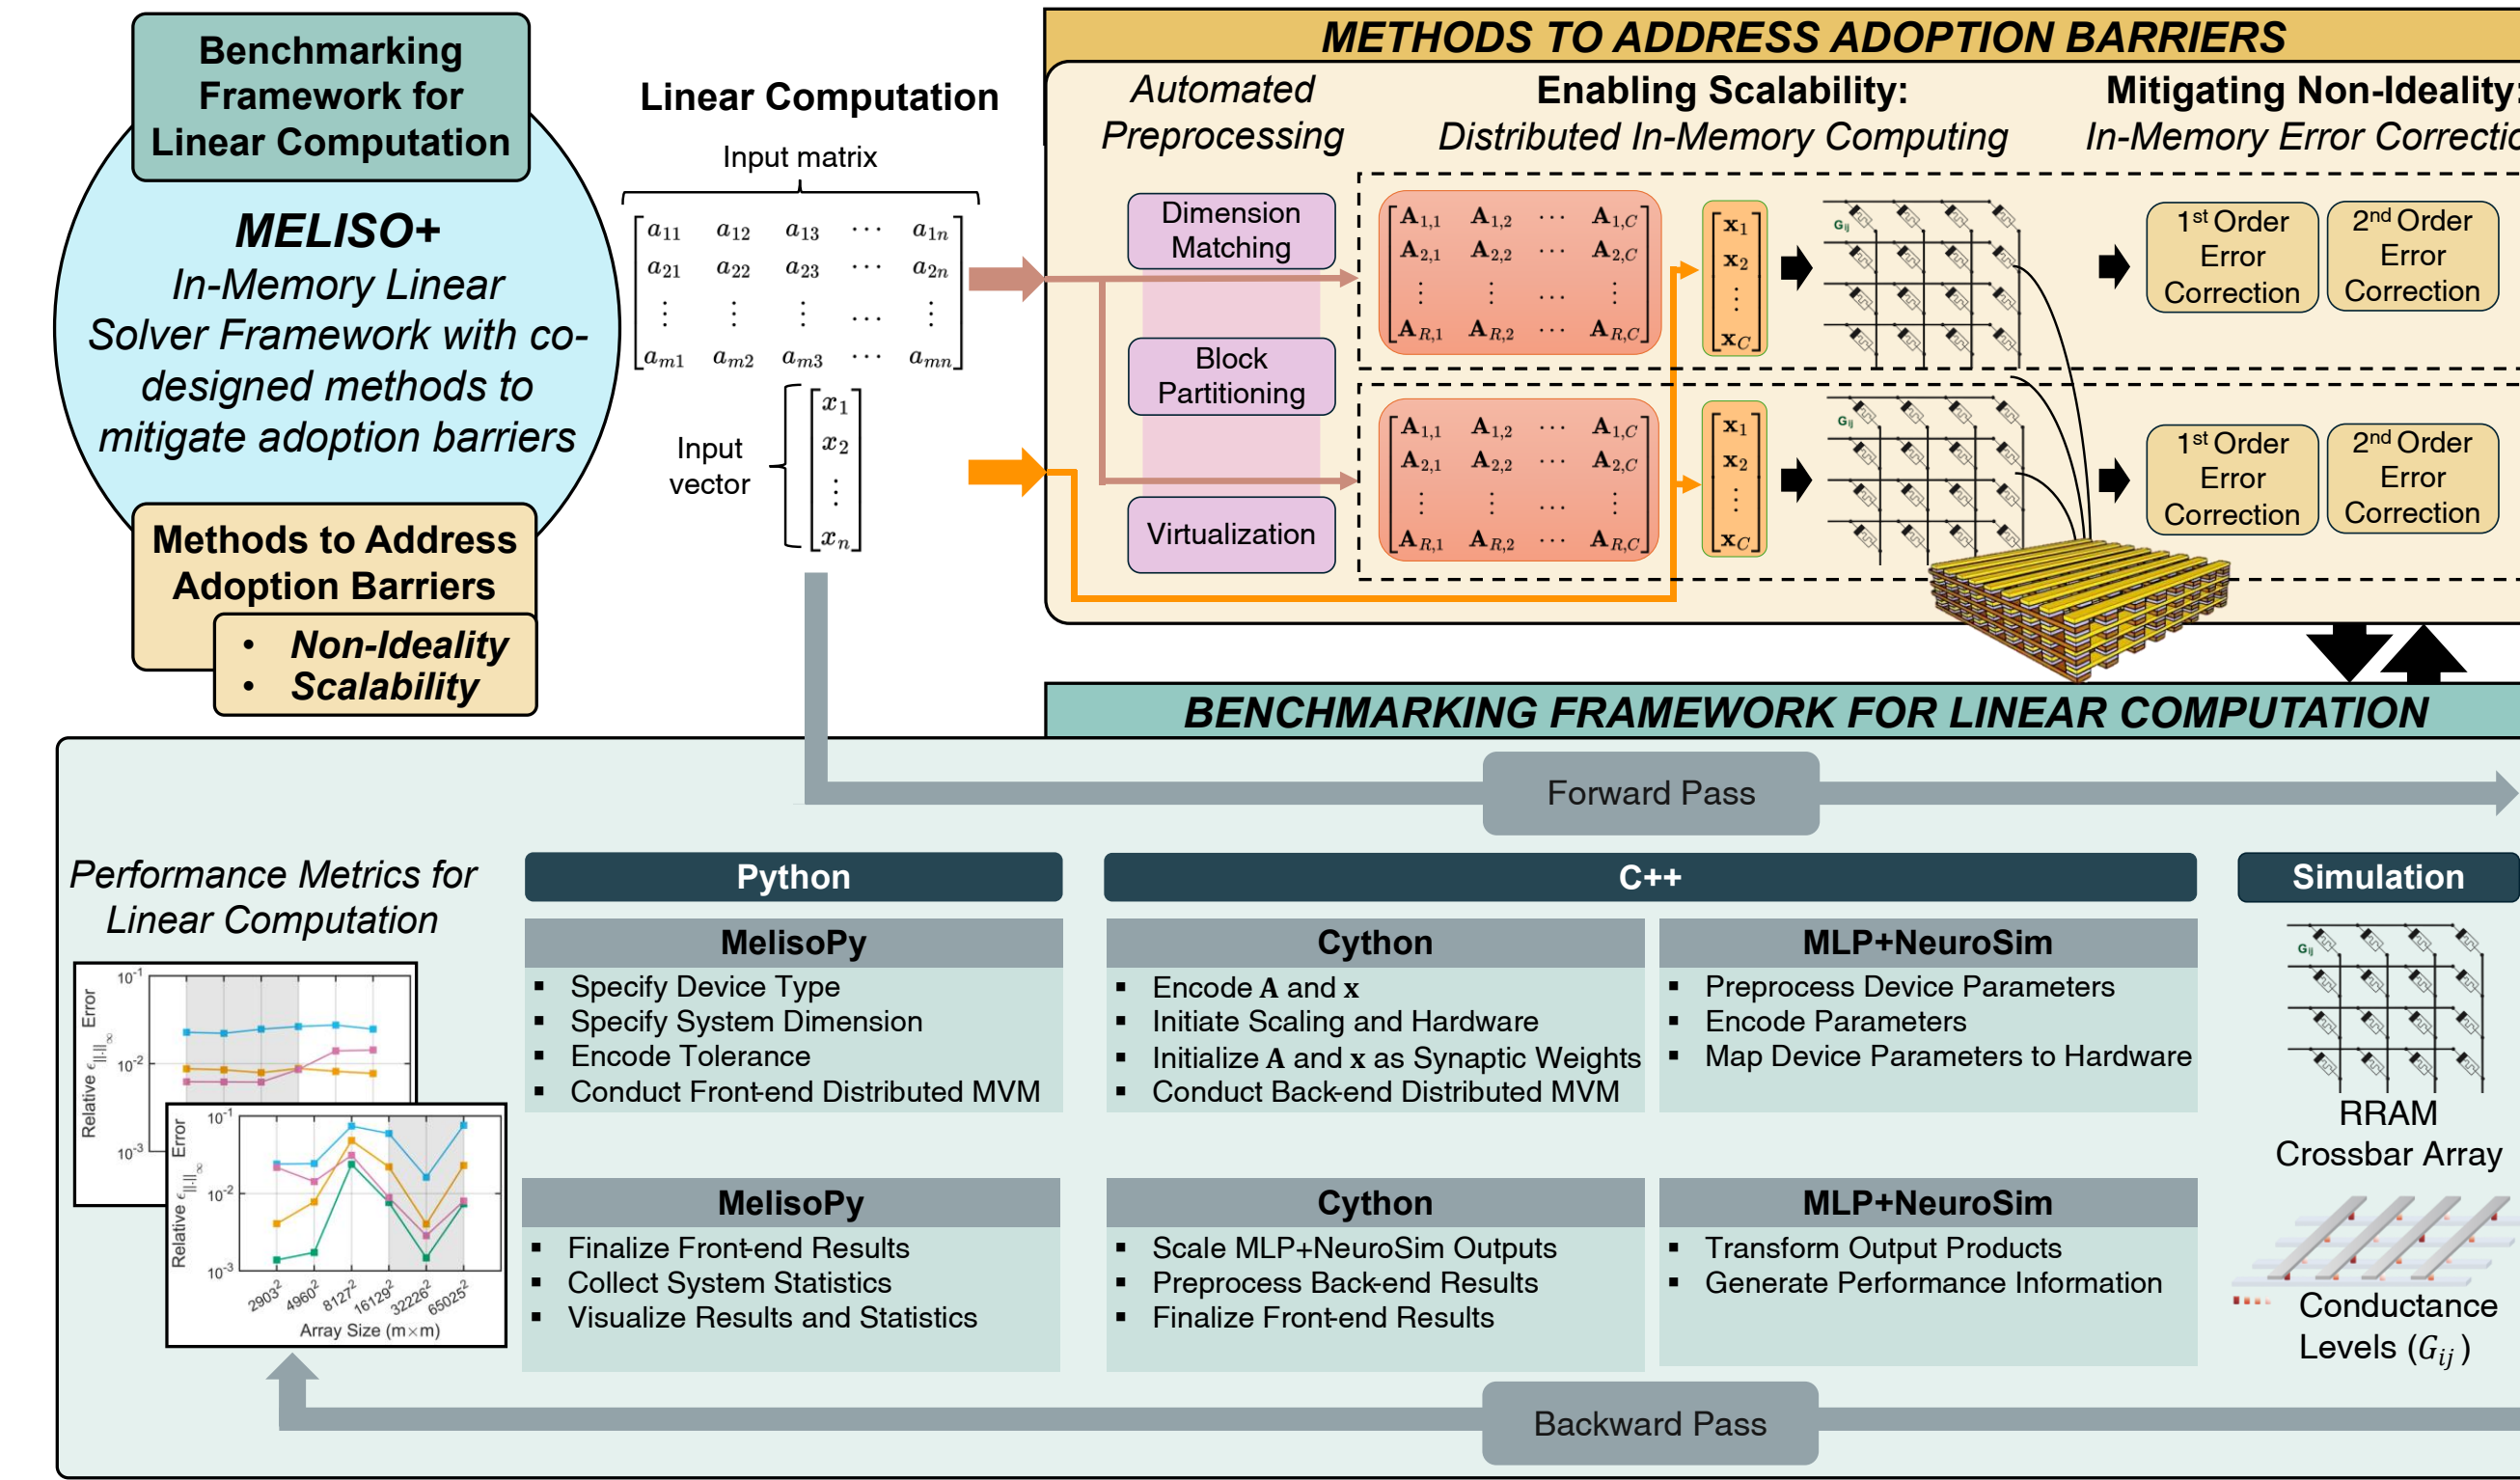


CHALLENGES — Each generation of solver trades one limitation for another

- Exact digital solvers:** Classical methods (e.g. Gurobi) rely on repeated **matrix factorizations**. They return accurate, certified solutions but incur prohibitive memory overhead at scale.
- GPU first-order methods:** Newer methods such as **PDHG** exploit the abundance of GPUs to avoid factorization altogether, yet remain bound by the **von Neumann bottleneck**.
- Analog in-memory computing:** Devices such as **RRAM** are **noisy and non-ideal**. Additionally, the **writing operations are expensive** in RRAMs, so the problem instances (a.k.a. the matrices) *must not* be re-encoded each iteration.

METHODOLOGY — The MELISO+ simulation framework

We first build **MELISO+**, an in-memory linear-solver framework that simulates **RRAMs** under realistic device physics to measure **accuracy, latency, and energy**. It establishes how to perform the core primitive (**matrix-vector multiplication**, abbreviated as **MVM**) reliably on RRAM crossbar arrays by addressing two adoption barriers: **scalability** and **device non-ideality**.



- Message Passing Interface (MPI)** framework is employed to distribute the problem instances to multiple crossbar arrays (MCAs).
- Device properties** (e.g., device-to-device and cycle-to-cycle variations) are simulated to observe their effects on MVM results.

METHODOLOGY — In-memory optimization with PDHG

We then configure MELISO+ to solve linear programs entirely in memory using the **PDHG** method. Recast as a saddle-point problem,

$$\min_{l \leq \mathbf{x} \leq u} \max_{\mathbf{y}} \mathbf{c}^\top \mathbf{x} + \langle \mathbf{A}\mathbf{x} - \mathbf{b}, \mathbf{y} \rangle,$$


each iteration touches the matrix *only* through matrix–vector products:

$$\begin{aligned} \mathbf{y}_{k+1} &= \Pi_{\mathbf{y}}[\mathbf{y}_k + \sigma(\mathbf{A}\bar{\mathbf{x}}_k - \mathbf{b})], \\ \bar{\mathbf{x}}_{k+1} &= \Pi_{[l,u]}[\mathbf{x}_k - \tau(\mathbf{c} + \mathbf{A}^\top \mathbf{y}_{k+1})], \\ \bar{\mathbf{x}}_{k+1} &= \bar{\mathbf{x}}_{k+1} + \theta(\bar{\mathbf{x}}_{k+1} - \bar{\mathbf{x}}_k). \end{aligned}$$

Here, the products $\mathbf{A}\bar{\mathbf{x}}_k$ and $\mathbf{A}^\top \mathbf{y}_{k+1}$ are the in-memory MVMs; Π are cheap projections, $\{\sigma, \tau\}$ is the dual-primal step size, and θ is the extrapolation term.


We introduce the following modifications to ensure feasible implementation on RRAM devices:

- Write-once encoding.** The problem is cast as a single *symmetric block matrix* $\mathbf{M} = \begin{bmatrix} \mathbf{0} & \mathbf{A} \\ \mathbf{A}^\top & \mathbf{0} \end{bmatrix}$ and programmed onto the crossbars **exactly once**; every iteration reuses this encoding, with only the input vectors changing.
- Distributed single accelerator.** A 4×4 tile of 64×64 crossbars (a 256×256 logical array) operates as one accelerator, eliminating inter-tile data movement.
- Accelerated convergence.** The PDHG iteration is augmented with **Nesterov momentum** for adaptive step-size rule, reducing the number of iterations required to reach a target tolerance.
- Robustness by design.** Stable operator-norm estimation (**Lanczos**) and **write-and-verify** programming mitigate device non-idealities, and we **prove convergence on noisy hardware** — no worse than a digital baseline up to a small, fixed error term.





EXPERIMENTS AND RESULTS

Benchmarked against **gpuPDL** (a GPU version of the same method), with **Gurobi** as ground truth, on real LPs from MIPLIB-2017. Both RRAM devices (**TaO_x-HfO_x** and **EpiRAM**) reach solution within 1% GPU/Gurobi accuracy (relative gap Δ_{rel}) at a fraction of the energy and latency.



Total energy and latency measurement scheme



CONCLUSION

- The **first PDHG linear-program solver** simulated on RRAM — a true **algorithm–hardware co-design** (write-once encoding + noise-robust iteration).
- Provable** convergence on noisy analog hardware; **orders-of-magnitude** energy and latency gains over a GPU.

REFERENCES & ACKNOWLEDGMENTS

- [1] Vo, Chowdhury, Ramanan *et al.* From GPUs to RRAMs: Distributed In-Memory Primal–Dual Hybrid Gradient Method for Solving Large-Scale Linear Optimization Problems. *SIAM PP* (2026). arXiv:2509.21137.
 - [2] Vo, Chowdhury, Ramanan *et al.* Harnessing the full potential of RRAMs through scalable and distributed in-memory computing with integrated error correction. *Nature Communications Engineering* (2026). arXiv:2508.13298
 - [3] Chowdhury *et al.* The Lynchpin of In-Memory Computing: A Benchmarking Framework for Vector-Matrix Multiplication in RRAMs. *ICONS* (2024). arXiv:2409.06140
 - [4] Lu *et al.* A GPU Implementation of Restarted Primal-Dual Hybrid Gradient for Linear Programming in Julia. arXiv:2311.12180
- Conducted at Oklahoma State University, Argonne National Laboratory, and Wayne State University. Supported in part by the National Science Foundation Secure and Trustworthy Cyberspace (SaTC) program under Award No. 2348411.